




TensorFlow

AN OPEN-SOURCE SOFTWARE LIBRARY FOR MACHINE INTELLIGENCE



What is tensorflow?

- Open-source library by Google for **numerical computation** (not just deep learning!)
- Build a **computational graph** with the mathematical operations you want to perform
- Core **backend** developed in **C++** for efficiency, supports multiple CPUs, multiple GPUs with use of **CUDA** for GPU acceleration
- Interface with the backend using APIs for **Python**, Java, Go, ...
- Multiplatform: Linux, MacOS, Windows, Android, iOS
- Not the only existing one: Torch (Facebook), Theano (U. of Montreal), ...






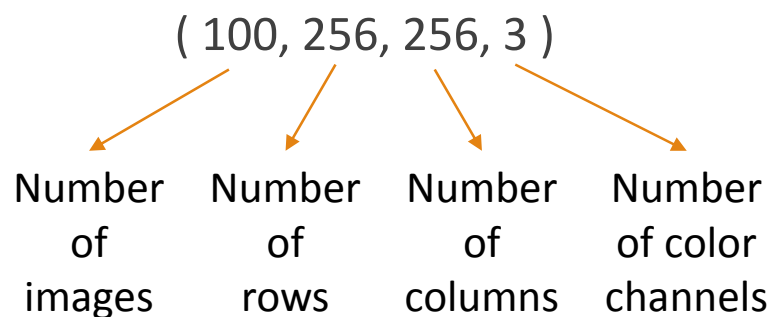


What is a tensor?

Tensor = multidimensional array

- 1D tensor: vector
- 2D tensor: matrix
- 3D tensor : RGB image

Example: a sequence of images as a 4D tensor



$\mathbf{x} \in \mathbb{R}^{4 \times 4 \times 4}$



Guide to the Python API

import tensorflow as tf

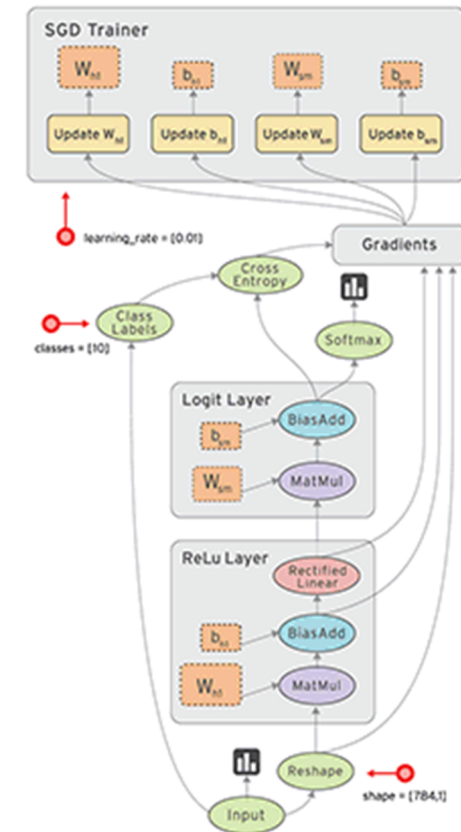
Import the tensorflow
Python module

Commonly used
shorthand



Guide to the Python API

1. Define a computational graph
2. Create a **Session** to run parts of the graph with some input data





Types of nodes

- **Constant**

➤ A tensor with constant value

```
c = tf.constant(4.0, dtype=tf.float32)
```

↓
value

↓
type:

tf.float32 (default)

tf.int8

tf.uint8

tf.int16

tf.int32

tf.bool

tf.string

...

TF documentation: class tf.DType



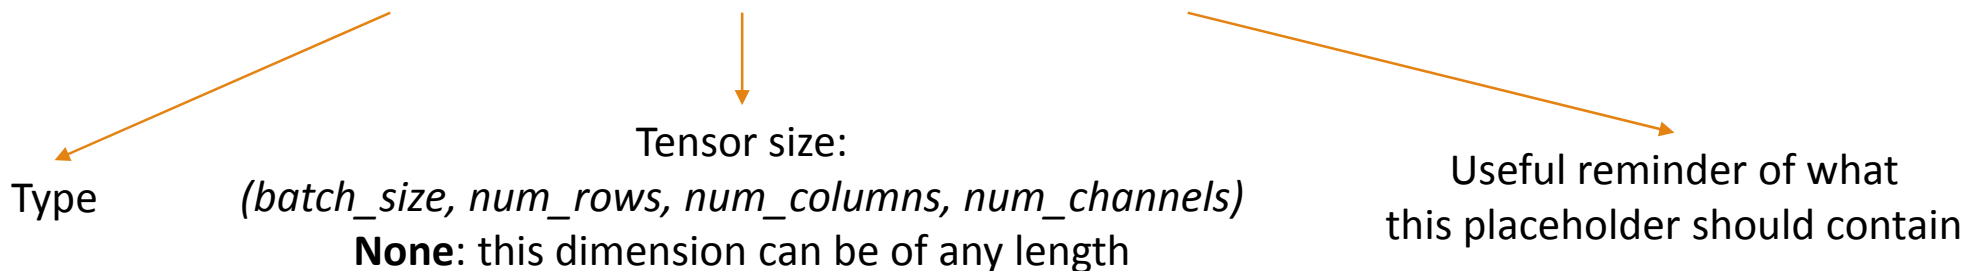
Types of nodes

- Placeholder

- Placeholders are used to *provide inputs*, i.e. the placeholder will be filled in with some value provided at runtime.

Example: define a placeholder that at runtime will hold a batch of images to be processed

```
x = tf.placeholder(tf.float32, shape=(None, 256, 256, 3), name='input_images')
```



Tip: defining the dimension of the placeholder corresponding to the number of samples in the batch as *None* allows you to run the same code for any batch size without modifications



Types of nodes

- **Variable**

➤ Variables are used to define the *trainable tensors* of the graph, e.g. the layer weights in a neural net.

```
W = tf.Variable(tf.truncated_normal(shape=(100,10), sdtdev=0.1), name='trainable_weights')
```



Initial value

`tf.truncated_normal`: random Gaussian (without extreme values) tensor

`tf.constant`: constant tensor

`tf.zeros`: constant tensor filled with 0s

`tf.ones`: constant tensor filled with 1s

...

TF Documentation:

https://www.tensorflow.org/api_guides/python/constant_op



Useful reminder of what this variable is



Mathematical operations

- Tensorflow provides functions to compute most mathematical operations. These functions call efficient implementations in the backend
 - For these functions Tensorflow knows both the
 - **forward** behavior: from inputs and internal parameters compute output
 - **backward** behavior: compute derivative of loss with respect to input and internal parameters
- This allows optimizers to apply the **chain rule** to compute derivatives across all the operations in the graph

```
h = tf.matmul(x, W) + b # tensor product and sum x*W+b  
h = x*W + b # same as before (*,+) are recognized by tf
```



A few useful operations

```
out = tf.nn.relu(input) # ReLU activation
```

```
out = tf.nn.sigmoid(input) # sigmoid activation
```

```
out = tf.nn.conv2d(x, W, strides=(1,1,1,1),padding='SAME') # 2D convolution
```

```
out = tf.nn.max_pool(input, ksize=(1,2,2,1), strides=(1,2,2,1), padding='SAME') # Max pooling
```

```
out = tf.nn.dropout(input, keep_prob) # Dropout, keep_prob is a placeholder
```

```
out = tf.equal(a, b) # Returns the boolean truth value of (x == y) element-wise
```

```
out = tf.cast(input, tf.float32) # cast input tensor to specified type
```

```
out = tf.argmax(input, axis=1) # Returns the index with the largest value across the specified dimension (axis) of a tensor
```

```
out = tf.reduce_mean(input, axis=0) # Computes the mean of elements across the specified dimension of a tensor. e.g.: if input has shape (100,10), axis=0, output has shape (1,10)
```

https://www.tensorflow.org/api_guides/python/ for many more

Note on multinomial logistic regression



- Multinomial logistic regression = logistic regression for **multiclass classification** problems
- Last layer of network must have
 - **Softmax** activation $s_i = \text{softmax}(h_i) = \frac{e^{h_i}}{\sum_j e^{h_j}}$
 - **Cross-entropy** loss $D(\mathbf{s}, \mathbf{y}) = -\sum_i y_i \log(s_i)$

Tensorflow allows you to do this but THIS IS NOT NUMERICALLY STABLE and will work poorly

```
s = tf.nn.softmax(h) # softmax  
loss = - tf.reduce_sum(y * tf.log(s), 1) # cross entropy
```

DO NOT DO THIS!

Instead, the softmax and cross-entropy are combined in a single **numerically-stable** operation:

```
loss = tf.nn.softmax_cross_entropy_with_logits(labels=y, logits=h) # numerically stable
```



Optimization

- Now that you defined all the computations and the loss function, you must specify if you want to **optimize some variable** and **how**
- The **train API** provides optimizers to minimize loss functions with respect to Variables

```
optimizer = tf.train.GradientDescentOptimizer(0.5) # 0.5 is the learning rate  
train_step = optimizer.minimize(loss)
```

train_step is a new node in the graph, executing it runs **one iteration** of gradient descent

Optionally:
optimizer.minimize(loss, var_list=...)
Provide a list of variable instead of optimizing over all of them

- tf.train has several functions such as more advanced optimizers (tf.train.AdamOptimizer, ...) and options such as learning rate decay



Sessions – Creating them

- You need a **Session** to run a graph or parts of it



Connection to
the C++ backend

```
sess = tf.InteractiveSession()
```

- *tf.InteractiveSession()* is the same as *tf.Session()* except it is less annoying because it automatically sets itself as default session



Sessions – Initialize variables

- Before training anything make sure your **variables are initialized** by calling

```
tf.global_variables_initializer().run()
```

- This operation will initialize all the variables using the **initialization functions or tensors** you specified when you declared them with `tf.Variable`



Sessions – Run

- Once you have a Session (and have initialized the Variables if you want to train them), you can **run parts of the graph**

```
sess.run(train_step, feed_dict={x: batch_xs, y: batch_ys})
```

↘ **feed_dict** is a **Python dictionary** (syntax is *key: value*) that takes tensors that will be put inside **placeholders**.

The key is the placeholder name.

The value is a tensor (this is an input, not part of the graph, it has numbers in it)

↓
This is the **node of the graph** we want to run: in this case is a training operation of the optimizer that will **run one iteration** of SGD



Sessions – Run

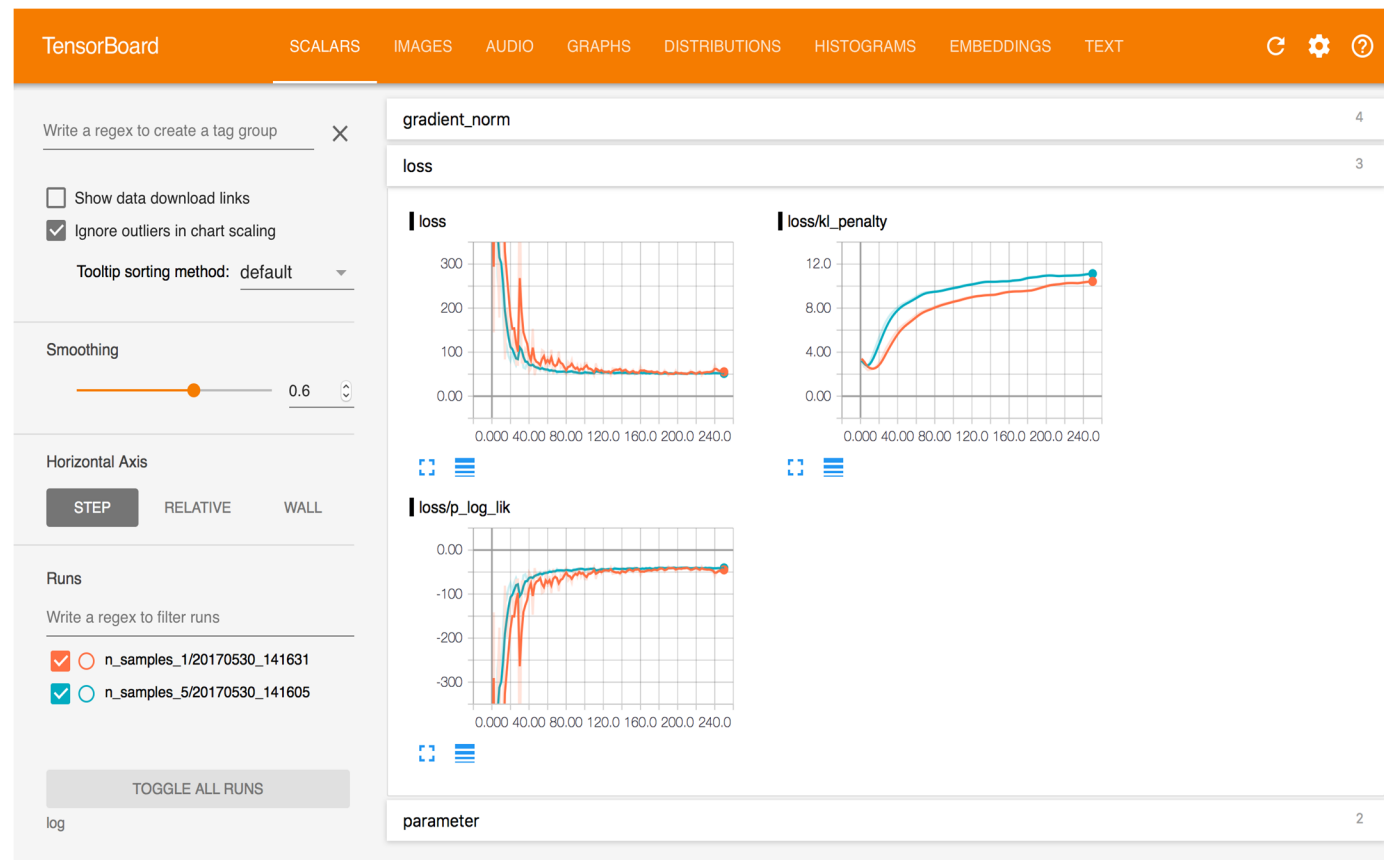
- You can **run any node** in the graph to get its **value** (it will perform a forward pass from the input to node you specified)
- Suppose you want to compute an accuracy value of your method

```
accuracy = ... # define the operation to compute accuracy as node in the graph  
accuracy_value = sess.run(accuracy, feed_dict={x: batch_xs, y: batch_ys}) # run the operation to get the value
```



Tensorboard

- Tensorboard is a **graphical interface** to visualize tensor values, statistics, the computational graph , ...
- **Very useful for debugging:** keep track of the loss function during a long training, visualize a histogram of weight values
- You must **write instructions** in your program to **save** information to an **Event File** that is read by Tensorboard for visualization





Tensorboard - Summaries

- When you have a **tensor** you want to keep track of, you attach a **tf.summary** to it
- There are different types of summaries depending on the tensor and visualization:

```
tf.summary.scalar('loss', loss) # a scalar, e.g. loss function or accuracy
```

```
tf.summary.histogram('weights_layer1', W1) # a histogram of all the values in the tensor, e.g. check that your weights are not all 0s!
```

```
tf.summary.image('input_images', x) # visualize a batch of images (batch_size, Nrows, Ncols, Nchannels)
```

- They are all nodes in the graph; once you defined all the **summary nodes** you want, you need to **run them** to get the actual values
- **Merge all the summaries** into one object for simplicity:

```
merged = tf.summary.merge_all()
```



Tensorboard - Event File

- All the data in summaries must be written into a file that is read by the Tensorboard program
- To do that we use **FileWriter** that takes the **merged summaries** and **serializes them to a file**

1. Instantiate the FileWriter object

```
train_writer = tf.summary.FileWriter(FLAGS.log_dir + '/train', sess.graph)
```

Logging directory where to write the file

Notice you need a session

2. Run the merged summaries to get the values

```
summary_train, _ = sess.run([merged, train_step], feed_dict={x: batch_xs, y: batch_ys}) # do one step of training and compute summaries
```

3. Write summaries

```
train_writer.add_summary(summary_train, i) # i is the iteration counter
```

Keras



- **Keras** is a high-level API written in Python to code neural networks more easily
- It runs **on top of Tensorflow** (or Theano)
- Keras can be used independently (without ever writing Tensorflow code) or can be integrated into the native Tensorflow flow



Independent Keras

Its own model compiling,
optimization, ...
Agnostic to backend (can
run on Theano without
code modifications)



tf.contrib.keras

Use tensorflow optimizers, Tensorboard, ...
WE WILL USE THIS


Note: from tensorflow 1.4, Keras is no longer
a contributor package but officially supported
as **tf.keras**



Keras - Models

- Everything works as in pure Tensorflow (placeholders, sessions, tensorboard)
- Except there are high-level definitions for many common neural networks layers so that you don't have to worry about variables
- A **Keras model is a sequence of layers**
 - **Sequential model:** linear stack of layers
 - **Functional API:** allows to defined arbitrarily complex arrangements of layers

Keras – Sequential model



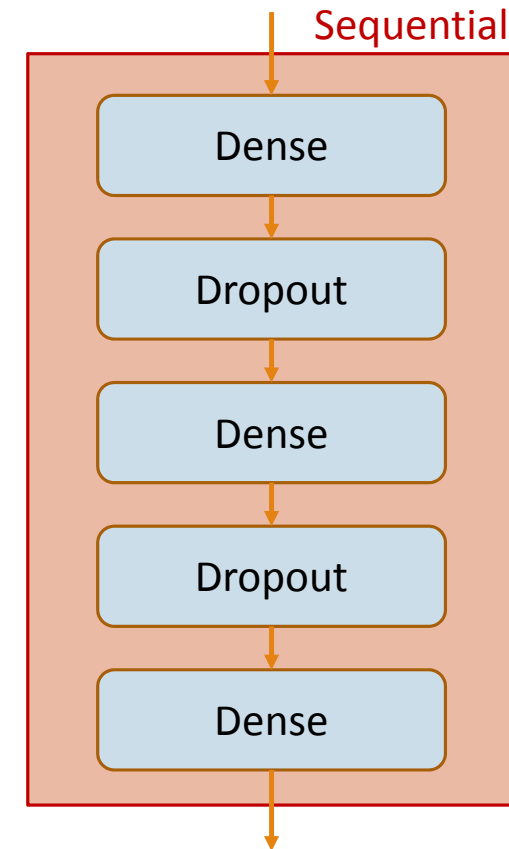
- **Linear stack of layers:** the output of a layer is the input of the next layer

```
def my_model():  
  
    model = models.Sequential()  
    model.add(layers.Dense(64, activation='relu', input_dim=20))  
    model.add(layers.Dropout(0.5))  
    model.add(layers.Dense(64, activation='relu'))  
    model.add(layers.Dropout(0.5))  
    model.add(layers.Dense(10))  
  
    return model
```


Create a
Sequential model

Add layers
with .add()

First layer only:
must specify input
dimension
(without batch size)



Keras – Layers, Initializers, Regularizers



- Visit <https://keras.io/layers/about-keras-layers/> for a comprehensive list of layers. Examples:

```
layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform', bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None, kernel_constraint=None, bias_constraint=None)
```

```
layers.Conv2D(filters, kernel_size, strides=(1, 1), padding='valid', data_format=None, dilation_rate=(1, 1), activation=None, use_bias=True, kernel_initializer='glorot_uniform', bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None, kernel_constraint=None, bias_constraint=None)
```


```
layers.Flatten()
```

- For the initializer parameters you can use the Keras initializer functions:

```
initializers.Zeros()  
initializers.Ones()  
initializers.RandomNormal(mean=0.0, stddev=0.05, seed=None)  
initializers.glorot_normal(seed=None)
```

- For the regularizers you can use:

```
regularizers.l1(0.01)  
regularizers.l2(0.01)  
regularizers.l1_l2(0.01)
```

Keras – Using a model

- Once you create your model you can use it as a function on any tensor in the graph

```
my_net = my_model()  
h = my_net(x) # x is a placeholder
```

- Keras models have extra methods and attributes. For example:

```
my_net.summary()
```

will print a summary of the model with
all the tensor sizes and
number of trainable parameters

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_1 (Dense)	(None, 64)	1344
dropout_1 (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 64)	4160
dropout_2 (Dropout)	(None, 64)	0
dense_3 (Dense)	(None, 10)	650
=====	=====	=====
Total params: 6,154		
Trainable params: 6,154		
Non-trainable params: 0		